

PROJECT

Train a Smartcab to Drive

A part of the Machine Learning Engineer Nanodegree Program

PROJECT REVIEW	CODE REVIEW 2	NOTES
<div>▼ agent.py 2</div> <pre> 1 import random 2 import math 3 from environment import Agent, Environment 4 from planner import RoutePlanner 5 from simulator import Simulator 6 7 class LearningAgent(Agent): 8 """ An agent that learns to drive in the Smartcab world. 9 This is the object you will be modifying. """ 10 11 def __init__(self, env, learning=False, epsilon=1.0, alpha=0.5): 12 super(LearningAgent, self).__init__(env) # Set the agent in the env: 13 self.planner = RoutePlanner(self.env, self) # Create a route planner 14 self.valid_actions = self.env.valid_actions # The set of valid actions 15 16 # Set parameters of the learning agent 17 self.learning = learning # Whether the agent is expected to learn 18 self.Q = dict() # Create a Q-table which will be a dictionary 19 self.epsilon = epsilon # Random exploration factor 20 self.alpha = alpha # Learning factor 21 22 ##### 23 ## TO DO ## 24 ##### 25 # Set any additional class parameters as needed 26 self.counter = 0 27 random.seed(161) #https://stackoverflow.com/questions/22639587/random-seed 28 29 def reset(self, destination=None, testing=False): 30 """ The reset function is called at the beginning of each trial. 31 'testing' is set to True if testing trials are being used 32 once training trials have completed. """ 33 34 # Select the destination as the new location to route to 35 self.planner.route_to(destination) 36 37 ##### 38 ## TO DO ## 39 ##### 40 # Update epsilon using a decay function of your choice 41 # Update additional class parameters as needed 42 # If 'testing' is True, set epsilon and alpha to 0 43 self.counter=self.counter+1 44 if testing: 45 self.epsilon = 0 46 self.alpha = 0 47 else: 48 #Default Q Learning Parameter 49 #self.epsilon=self.epsilon-0.05 50 #Improved Q Learning agent 51 #self.epsilon=0.9**self.counter 52 #self.epsilon=1/(self.counter**3) 53 #self.epsilon=math.exp(-0.01*self.counter) 54 self.epsilon=math.cos(-0.001*self.counter) 55 56 return None 57 58 def build_state(self): 59 """ The build_state function is called when the agent requests data from 60 environment. The next waypoint, the intersection inputs, and the data 61 are all features available to the agent. """ 62 63 # Collect data about the environment </pre>		

```

64     waypoint = self.planner.next_waypoint() # The next waypoint
65     inputs = self.env.sense(self)           # Visual input - intersection l
66     deadline = self.env.get_deadline(self)   # Remaining deadline
67
68     #####
69     ## TO DO ##
70     #####
71
72     # NOTE : you are not allowed to engineer eatures outside of the inputs
73     # Because the aim of this project is to teach Reinforcement Learning, v
74     # constraints in order for you to learn how to adjust epsilon and alpha
75     # With the hand-engineered features, this learning process gets entirel
76     state=None
77     # Set 'state' as a tuple of relevant data for the agent
78     state = waypoint, inputs['light'], inputs['left'], inputs['right'], in
79
80     return state
81
82
83     def get_maxQ(self, state):
84         """ The get_max_Q function is called when the agent is asked to find th
85             maximum Q-value of all actions based on the 'state' the smartcab is
86
87         #####
88         ## TO DO ##
89         #####
90         # Calculate the maximum Q-value of all actions for a given state
91         maxQ=None
92         maxQ = self.Q[state][max(self.Q[state], key=lambda key: self.Q[state][k
93
94         return maxQ
95
96
97     def createQ(self, state):
98         """ The createQ function is called when a state is generated by the age
99
100        #####
101        ## TO DO ##
102        #####
103        # When learning, check if the 'state' is not in the Q-table
104        # If it is not, create a new dictionary for that state
105        # Then, for each action available, set the initial Q-value to 0.0
106        if self.learning:
107            if state not in self.Q:
108                self.Q[state]={None:0.0, 'forward':0.0, 'left':0.0, 'right':0.0
109        return
110
111
112     def choose_action(self, state):
113         """ The choose_action function is called when the agent is asked to cho
114             which action to take, based on the 'state' the smartcab is in. """
115
116         # Set the agent state and default action
117         self.state = state
118         self.next_waypoint = self.planner.next_waypoint()
119         action = None
120
121         #####
122         ## TO DO ##
123         #####
124         # When not learning, choose a random action
125         # When learning, choose a random action with 'epsilon' probability
126         # Otherwise, choose an action with the highest Q-value for the current
127         # Be sure that when choosing an action with highest Q-value that you re
128         if not self.learning:
129             action = random.choice(self.valid_actions)
130         else:
131             if self.epsilon >= random.random():
132                 action = random.choice(self.valid_actions)
133             else:
134                 # Following piece of code implemented after 1st reviewer commen
135                 max_random_valid_actions = []
136                 max_Q_Value = self.Q[state][max(self.Q[state], key=lambda key:
137                 for act in self.Q[state]:
138                     if max_Q_Value == self.Q[state][act]:
139                         max_random_valid_actions.append(act)
140                 action = random.choice(max_random_valid_actions)
141
142         # Following piece of code implemented after 1st reviewer commen
143         #action = max(self.Q[state], key=lambda key: self.Q[state][key]
144         return action
145
146

```

AWESOME

Excellent implementation of a tie-breaker function between best actions! For this part of the code, you can use:

```

action = random.choice([k for (k, v) in Q_state.items() if v == max_Q_Value])

```

```

146 def learn(self, state, action, reward):
147     """ The learn function is called after the agent completes an action and
148         receives a reward. This function does not consider future rewards
149         when conducting learning. """
150
151     #####
152     ## TO DO ##
153     #####
154     # When learning, implement the value iteration update rule
155     # Use only the learning rate 'alpha' (do not use the discount factor)
156     if self.learning:
157         self.Q[state][action] = self.Q[state][action] + self.alpha*(reward-

```

AWESOME

Great job, your Bellman equation uses only current rewards.

```

158         return
159
160     def update(self):
161         """ The update function is called when a time step is completed in the
162             environment for a given trial. This function will build the agent
163             state, choose an action, receive a reward, and learn if enabled. """
164
165         state = self.build_state() # Get current state
166         self.createQ(state) # Create 'state' in Q-table
167         action = self.choose_action(state) # Choose an action
168         reward = self.env.act(self, action) # Receive a reward
169         self.learn(state, action, reward) # Q-learn
170
171         return
172
173     def run():
174         """ Driving function for running the simulation.
175             Press ESC to close the simulation, or [SPACE] to pause the simulation.
176
177             #####
178             # Create the environment
179             # Flags:
180             #   verbose - set to True to display additional output from the simulation
181             #   num_dummies - discrete number of dummy agents in the environment, default is 1
182             #   grid_size - discrete number of intersections (columns, rows), default is 10x10
183             env = Environment(verbose=True)
184
185             #####
186             # Create the driving agent
187             # Flags:
188             #   learning - set to True to force the driving agent to use Q-learning
189             #   * epsilon - continuous value for the exploration factor, default is 1
190             #   * alpha - continuous value for the learning rate, default is 0.5
191             agent = env.create_agent(LearningAgent, learning=True, alpha=0.95, epsilon=1)
192
193             #####
194             # Follow the driving agent
195             # Flags:
196             #   enforce_deadline - set to True to enforce a deadline metric
197             env.set_primary_agent(agent, enforce_deadline=True)
198
199             #####
200             # Create the simulation
201             # Flags:
202             #   update_delay - continuous time (in seconds) between actions, default is 1
203             #   display - set to False to disable the GUI if PyGame is enabled
204             #   log_metrics - set to True to log trial and simulation results to /logs
205             #   optimized - set to True to change the default log file name
206             sim = Simulator(env, update_delay=0.01, log_metrics=True, display=False, optimized=False)
207
208             #####
209             # Run the simulator
210             # Flags:
211             #   tolerance - epsilon tolerance before beginning testing, default is 0.01
212             #   n_test - discrete number of testing trials to perform, default is 10
213             sim.run(n_test=40, tolerance=0.0001)
214
215         if __name__ == '__main__':
216             run()

```

► logs/sim_improved-learning.txt

► logs/sim_default-learning.txt

RETURN TO PATH

[Student FAQ](#)